



Appliance Energy Prediction

Multivariate Time-Series Forecasting



Gamodya Rajasooriya



21/07/2026



AI Intern Technical Assessment

Table of Contents

Table of Contents.....	2
1. Introduction	3
2. Data Insights.....	3
Key findings	6
3. Preprocessing.....	7
3.1 Missing Values	7
3.2 Outlier Detection and Treatment	7
3.3 Target Transformation	8
3.4 Data Scaling.....	9
3.5 Train/Test Split	9
4. Feature Engineering.....	10
Features created	10
Feature Selection	11
5. Model Design	11
Baseline models.....	11
Deep learning models	11
Architecture experimentation (not fixed guesses)	11
Final architecture design.....	12
6. Results	12
Full model comparison	14
Key observations	14
7. Model Optimization.....	14
XGBoost tuning	15
LSTM tuning	15
Cross-cutting finding	15
8. Challenges and Solutions.....	15
9. Conclusion	16
10. References	16

1. Introduction

This project addresses the task of predicting household appliance energy consumption (Wh) using the multivariate Appliance Energy Prediction Dataset (Candanedo et al., 2017): 19,735 records collected at 10-minute intervals over 137 days from a low-energy house in Stambruges, Belgium, covering indoor temperature/humidity across 9 rooms, outdoor weather conditions, and the target variable, Appliances (Wh).

Rather than following a standard preprocessing-to-modeling pipeline by default, this project validated every major methodological decision through experiments. Decisions regarding outlier handling, target transformation, scaling method, train/test ratio, sequence window length, optimizer, dropout rate, and hyperparameter selection were treated as empirical questions rather than assumptions. This report documents the complete experimental process, including what was tested, the evidence obtained from each experiment, and the rationale behind every final decision.

To ensure transparency, each stage of the modelling process was supported by measurable results:

- **Preprocessing:** The original usage spikes were preserved, while the Box-Cox transformation ($\lambda = -0.531$) reduced the target skewness from 3.386 to -0.052, creating a more balanced distribution for modelling.
- **Feature Engineering:** Feature analysis showed that recent usage patterns, such as 1-hour rolling averages and lag values, contributed around 75% of the predictive power. This highlights that previous consumption behavior is more influential than individual sensor readings.
- **Deep Learning Tuning:** Experiments found that a 2-hour input window (12-time steps) with the Adam optimizer provided the most stable performance for sequence-based models.
- **Model Performance:** Both tuned models outperformed the Naive persistence baseline (MAE: 26.25 Wh, R^2 : 0.450). The tuned LSTM achieved a slightly better result with an MAE of 26.17 Wh and R^2 of 0.493, while tuned XGBoost delivered the strongest performance, achieving an MAE of 17.91 Wh and explaining 74.4% of the variance ($R^2 = 0.744$).

2. Data Insights

Figure 1: Timeline Analysis of Hourly Average Appliance Energy Consumption (Wh)




Figure 1 shows the daily energy consumption pattern, with two peaks observed. A smaller increase occurs around 11:00 AM, while the main peak appears in the evening (18:00–20:00), reaching approximately 190 Wh.

Figure 2: Cross-Sectional Bar Chart of Average Weekly Energy Load by Day of Week




Figure 2 highlights the weekly energy consumption pattern. Monday records the highest average usage at 111.5 Wh, while Tuesday to Thursday show the lowest consumption levels (87–90 Wh), indicating that energy demand does not simply increase during weekends.

Figure 3: Macro Temporal Trend Analysis Across the 137-Day Observational Horizon




Figure 3 shows the long-term energy consumption trend over 137 days. The moving average remains relatively stable from January to May, indicating no significant seasonal changes and suggesting that short-term patterns contribute more to consumption variations.

FIGURE 4: Correlation heatmap of all features




Figure 4 shows the correlation patterns between input features and energy consumption. Weather variables have weak linear relationships with the target ($r < 0.20$), indicating that single measurements alone have limited predictive power and supporting the use of rolling time-based features.

Figure 5: Target Variable (Appliances) Distribution Boxplot and Density Analysis




Figure 5 shows the right-skewed distribution of the original energy consumption values. The long tail reaching around 1,080 Wh indicates occasional high-demand periods, highlighting the need for variance stabilization before applying linear models.

Figure 6: Autocorrelation (ACF) and Partial Autocorrelation (PACF) Target Diagnostics




Figure 6 shows the temporal dependency of household energy usage. The PACF results show a strong relationship with the previous 1–3-time steps, supporting the use of early lag features to capture short-term consumption patterns.

Figure 7: Distribution Plot of target variable




Figure 7 shows that appliance energy consumption is highly right skewed, with most values concentrated at lower consumption levels.

Key findings

Data integrity: Zero missing values confirmed across all 29 original columns; timestamp intervals verified as consistently 10 minutes apart with no gaps.

Target distribution: Appliances is heavily right-skewed (skewness = 3.39), ranging 10–1080 Wh with mean 98 Wh but median only 60 Wh - motivating a transformation before modeling.

Weak individual correlations: No single feature correlated above ~0.2 with the target (strongest: lights at 0.197, RH_out at 0.15). This was a pivotal early finding - it indicated that instantaneous sensor readings alone would be insufficient predictors, motivating a shift toward temporal (lag/rolling) feature engineering.

Temporal patterns: A clear two-peak daily cycle emerged - a smaller mid-morning peak (~11 AM) and a larger evening peak (~6–8 PM, ~190 Wh average). By day of week, Monday showed the highest average consumption (111.5 Wh) - higher than either weekend day - contradicting the common assumption that weekends alone drive higher usage; Tuesday–Thursday were lowest (87–90 Wh).

Weather correlation tested, not assumed: Daily-average outdoor temperature showed negligible correlation with daily-average Appliances usage ($r = 0.007$). This makes sense on reflection: the Appliances variable excludes HVAC/heating systems, so outdoor weather would not be expected to directly drive it.

rv1/rv2 identified: These two columns were confirmed identical to each other and are documented in the original paper as deliberately injected random noise variables, included to test whether feature selection methods correctly identify and discard non-predictive features.

Multicollinearity: Severe correlation among indoor temperature sensors T1–T9 was visible in the heatmap, flagging a risk for linear models and motivating an aggregated indoor_temp_avg feature.

ACF/PACF: PACF showed a sharp cutoff after lag 1–2, with small residual significance fading by approximately lag 10–15; ACF decayed slowly across all lags tested. This directly informed lag feature selection (Section 4) and sequence window length (Section 5).

3. Preprocessing

3.1 Missing Values

Zero missing values were confirmed via `.isnull().sum()`. No imputation was required; had missing values existed, time-aware interpolation would have been preferred over mean imputation to preserve sequential trend integrity.

3.2 Outlier Detection and Treatment

IQR-based outlier detection flagged a meaningful proportion of high-usage points as statistical outliers. Four treatment strategies were empirically compared using a Linear Regression sanity check (MAE, Wh).

Approach	MAE (Wh)
Keep outliers, raw target	52.29
Keep outliers, log1p-transformed target	42.45
Remove outliers (IQR)	16.39
Cap/winsorize at IQR bound	26.59

Decision: keep outliers and apply a log-transform to the target. Four outlier treatment strategies were evaluated empirically using a Linear Regression baseline. Although removing outliers produced the lowest MAE (16.39 Wh), this result was considered misleading because it eliminated approximately 8–10% of the test data, including genuine high-consumption events, making the prediction task artificially easier. Similarly, capping (winsorizing) preserved all observations but distorted the true magnitude of high-usage events by flattening valid signal. In contrast, applying a log-transform retained all data points, reduced the target skewness from 3.39 toward a more symmetric distribution, and achieved the best performance (42.45 Wh MAE) among approaches that preserved the complete dataset. Therefore, retaining outliers and applying a log-transform was selected as the most realistic and effective preprocessing strategy for real-world deployment.

Figure 8: Empirical Test MAE Variance Across Four Outlier Mitigation Strategies




Figure 8 compares different outlier handling strategies. Removing outliers reduces error artificially but loses important data patterns, while Log/Box-Cox transformation preserves the full timeline and stabilizes variance for better modelling.

3.3 Target Transformation

Three transforms were compared by resulting skewness and predictive MAE.

Transform	Resulting Skew	MAE (Wh)
Raw (none)	3.386	52.29
Log1p	1.190	42.45
Box-Cox ($\lambda = -0.531$)	-0.052	41.49
Yeo-Johnson	-0.046	(not separately tested; near-identical to Box-Cox)

Decision: Box-Cox transformation ($\lambda = -0.531$). Four target transformation methods were evaluated using skewness reduction and predictive performance. Box-Cox reduced the target skewness from **3.386** to **-0.052**, substantially outperforming the log transform (skew = **1.190**) while also achieving the lowest MAE (**41.49 Wh** vs. **42.45 Wh** for log1p). Since the Appliances target contains only positive values, the Box-Cox assumption of strictly positive data is fully satisfied, making Yeo-Johnson's additional flexibility for zero or negative values unnecessary. Therefore, Box-Cox was selected as the final target transformation, and the fitted λ value was retained for inverse transformation during model evaluation.

FIGURE 9: Target transform comparison chart




Figure 9 shows the effect of mathematical transformations on the target distribution. The Box-Cox transformation successfully reduces skewness from 3.386 to -0.052, creating a more balanced distribution and improving suitability for modelling.

3.4 Data Scaling

Scaler selection was first evaluated using a Linear Regression baseline. All three scaling methods (RobustScaler, StandardScaler, and MinMaxScaler) produced identical performance (MAE = 63.12), as Ordinary Least Squares regression is mathematically unaffected by linear feature scaling. Therefore, RobustScaler was initially selected because it is less sensitive to retained outliers, using the median and interquartile range instead of the mean and standard deviation. However, a later experiment showed that neural networks responded differently to feature scaling. For the LSTM model, MinMaxScaler achieved the best performance (MAE = 26.15), outperforming RobustScaler (MAE = 28.91), while StandardScaler resulted in unstable training ($R^2 = -0.99$). Consequently, RobustScaler was retained for the tree-based models, whereas MinMaxScaler was adopted for the final LSTM model.

3.5 Train/Test Split

Three ratios were tested via Linear Regression (raw target, to isolate this variable from Box-Cox inversion sensitivity).

Split	MAE (Wh)	Test rows
70/30	69.66	5,921
80/20	63.12	3,947
90/10	52.31	1,974

Decision: 80/20 chronological train-test split. Although MAE consistently decreased as the training set became larger, this improvement was attributed to the progressively smaller test set rather than better model generalization. An 80/20 split provided the best balance between sufficient training data and a large, representative test set (approximately 3,947 observations, or 27 days) for reliable evaluation. To preserve the temporal order of the data and prevent data leakage, the earliest 80% of observations were used for training and the remaining 20% for testing. All preprocessing steps, including scaling and target transformation, were fitted exclusively on the training data before being applied to the test set.

4. Feature Engineering

Figure 10: Feature Importance Pareto Bar Chart of Top 15 Engineered Predictors




Figure 10 shows that engineered temporal features provide the strongest predictive information. Recent energy patterns, including rolling averages and lag values, capture most of the variance, indicating that historical load behavior is more influential than static weather features.

Features created

Time-based: hour, day-of-week, month, is_weekend

Cyclic encoding: sine/cosine transforms of hour and day-of-week, preventing artificial discontinuity between e.g. 23:00 and 00:00

Domain-specific: is_holiday flag (Belgian public holidays within the Jan–May 2016 window)

Lag features: Appliances at lags 1, 2, 3, 6, and 10 steps (10–100 minutes prior), selected based on the PACF plot's sharp cutoff and fading tail (Section 2)

Rolling statistics: 1-hour and 3-hour rolling mean and standard deviation of Appliances, motivated by the slowly decaying ACF, capturing longer trend momentum that discrete lags alone would not

Interaction feature: indoor-outdoor thermal gradient ($\text{indoor_temp_avg} - T_{\text{out}}$), physically motivated by HVAC-driven energy use, though later found to have negligible importance, consistent with the near-zero weather correlation found in EDA

Multicollinearity mitigation: indoor_temp_avg and indoor_temp_std, aggregating the 9 correlated indoor temperature sensors

Feature Selection

Random Forest feature importance was used for selection. Results revealed a striking concentration of predictive power.

Feature	Importance
roll mean 1h	60.4%
Appliances lag 1	14.9%
roll std 1h	4.6%
Appliances lag 3	3.2%
(all remaining 41 features combined)	~17%

Recent usage history, including rolling averages and immediate lag features, contributes around 75% of the model's predictive power. This strongly supports the earlier EDA finding that individual weather and sensor measurements provide limited standalone information about appliance energy consumption. The feature importance analysis also showed that weather and sensor variables each contributed less than 1% of the overall importance. Even the engineered thermal_gradient feature had only around 0.5% importance, which aligns with its near-zero correlation observed during EDA.

Interestingly, the engineered indoor_temp_avg feature did not significantly outperform the individual T1–T9 temperature sensors as initially expected. This suggests that the Random Forest model is already able to handle the existing multicollinearity among temperature sensors and extract useful patterns without requiring manual aggregation.

5. Model Design

Baseline models

Naive persistence, Linear Regression, Random Forest, and XGBoost were trained as benchmarks.

Deep learning models

Three architectures were designed and compared: LSTM, GRU, and a CNN-GRU hybrid (1D-CNN layer for local pattern extraction feeding into a GRU for longer-term sequence modeling, following Pan et al.'s hybrid approach).

Architecture experimentation (not fixed guesses)

Three ablation studies were conducted on the GRU architecture before finalizing settings:

Sequence window length (1hr / 2hr / 4hr of history):

Window	MAE	RMSE	MAPE	R ²
6	34.77	81.16	24.41%	0.162
12	29.41	72.90	22.14%	0.311
24	30.67	72.23	25.27%	0.324

Window=12 (2 hours) selected - wins MAE/MAPE outright; consistent with PACF's fading tail around lag 10–15.

Optimizer (Adam vs RMSprop):

Optimizer	MAE	RMSE	MAPE	R ²
Adam	29.52	74.11	21.63%	0.287
RMSprop	31.30	73.17	25.41%	0.305

Adam selected - wins the more interpretable metrics; combines momentum with adaptive per-parameter learning rates, suited to noisy, mixed-scale sensor data.

Dropout rate (0.0 / 0.2 / 0.3 / 0.5):

Dropout	MAE	RMSE	MAPE	R ²
0.0	29.38	69.18	24.15%	0.379
0.2	46.49	94.98	34.69%	-0.170
0.3	29.89	72.62	23.61%	0.316
0.5	37.03	88.06	24.21%	-0.006

Dropout=0.0 selected. Contrary to the common assumption that regularization improves generalization, no dropout performed best across most metrics - attributed to the network's relatively modest capacity and training set size (~15.8k rows), insufficient for dropout to help without discarding genuine signal.

Final architecture design

- Loss function: Huber loss (robust to the retained outlier spikes, consistent with prior experience on time-series regression)
- Activation: ReLU (dense layers), tanh (default LSTM/GRU gate activation)
- Layers: 2-layer LSTM/GRU (64→32 units), tapered structure; CNN-GRU uses a Conv1D (64, kernel=2) feature extractor followed by a single GRU (32) layer
- Regularization: no dropout (per ablation finding above); early stopping (patience=10, monitoring validation loss) used as the primary overfitting safeguard

6. Results

FIGURE 11: All-models comparison bar chart - MAE and R²




Figure 11 compares the performance of all tested models. Tuned XGBoost achieved the best results, with the lowest error of 17.91 Wh and explaining 74.4% of the variation in unseen data.

FIGURE 12: XGBoost predicted vs actual plot




Figure 12 shows the relationship between actual and predicted energy values. The model closely follows the 45-degree line for normal consumption levels, while higher energy surges show some underestimation due to extreme load variations.

FIGURE 13: XGBoost residual distribution and residuals-vs-predicted scatter




Figure 13 shows the error behavior of the best-performing model. The residual pattern indicates typical time-series variance changes, while the error distribution remains centered around zero, confirming reliable prediction performance.

Figure 14: Deep Learning Training Validation Loss Convergence Curves.




Figure 14 shows the training and validation performance of the three sequence models. The validation loss remains stable, indicating effective learning without overfitting due to the use of early stopping.

Full model comparison

Model	MAE (Wh)	RMSE	MAPE	R ²
Linear Regression	39.00	131.88	26.40%	-1.214
GRU	32.14	82.20	25.68%	0.123
CNN-GRU	31.49	73.43	25.24%	0.300
Naive (persistence)	26.25	65.75	21.90%	0.450
LSTM (final, tuned)	26.17	62.52	22.06%	0.493
Random Forest	21.47	52.67	16.50%	0.647
XGBoost (tuned) - BEST	17.91	44.86	14.61%	0.744

Key observations

XGBoost is the clear overall winner across every metric.

Linear Regression failed severely ($R^2 = -1.214$). Diagnosis: 40 of 3,930 predictions (~1%) hit the safety clip ceiling during Box-Cox inversion, and RMSE was 3.4× larger than MAE - the signature of a small number of catastrophic mispredictions. Root cause traced to unstable coefficients under the multicollinearity identified in EDA among T1–T9 sensors, combined with Box-Cox inversion sensitivity under a negative lambda.

LSTM outperformed GRU and the CNN-GRU hybrid among deep learning models, contradicting the literature-based hypothesis (Pan et al.) that a CNN-hybrid would outperform single-cell-type architectures - a reminder that findings from one dataset do not automatically transfer.

Tree-based models outperformed deep learning overall - a well-documented pattern for tabular/structured time-series problems at this data scale (~16k training rows), where deep architecture typically requires substantially more data to realize their representational advantages.

Residual analysis on the final XGBoost model revealed heteroscedasticity: error variance was small at low predicted values, widening substantially at higher predictions, with a systematic tendency to underestimate the sharpest consumption spikes. This is a genuine, reportable limitation - relevant if this model were deployed for peak-demand-sensitive applications such as grid capacity planning.

Benchmark comparison: Prior deep learning benchmarking studies on this identical dataset (Al-Obeidat et al., 2022; Candanedo et al., 2017) reported sequential baseline figures reaching **GRU $R^2=0.62$** . In our setup, **our tuned XGBoost model achieved $R^2=0.744$** (and test MAE=17.91 Wh), significantly outperforming those benchmark sequential figures. This confirms that the comprehensive rolling-average and lag feature engineering applied in our pipeline extracted substantial mathematical signal that traditional deep recurrent setups failed to capture natively.

7. Model Optimization

Figure 15: Hyperparameter Optimization Impact Matrix (Before vs. After Optuna Tuning)




Figure 15 shows the impact of hyperparameter tuning on model performance. XGBoost achieved a reduction in error from 19.29 Wh to 17.91 Wh, while the LSTM model showed limited improvement due to lower sensitivity to parameter changes.

Optuna (Bayesian/TPE-based search) was used for hyperparameter tuning, with TimeSeriesSplit (walk-forward) cross-validation to preserve temporal integrity - avoiding standard K-fold, which would leak future information into training.

XGBoost tuning

- v1 (30 trials; `n_estimators`, `max_depth`, `learning_rate`, `subsample`, `colsample_bytree`, `min_child_weight`): MAE improved 19.29 → 17.91, R^2 0.697 → 0.744.
- v2 (60 trials; added `gamma`, `reg_alpha`, `reg_lambda` regularization terms): MAE worsened to 23.14, R^2 dropped to 0.527. Diagnosis: added regularization over-constrained the model, causing underfitting despite a reasonable internal CV score. v1 retained as final.

LSTM tuning

Multiple tuning rounds were conducted:

- v1 (20 trials; `units`, `learning_rate`, `batch_size`): MAE 28.28 → 29.27 (slightly worse), but RMSE and R^2 improved (68.64→67.02, 0.389→0.417).
- v2 (25 trials; jointly tuning `window`, `layers`, `units`, `lr`, `batch_size`): mixed result vs v1 (MAE 28.93 vs 29.27, but R^2 0.365 vs 0.417) — the added architectural complexity (2 layers, 156 units) did not yield a clear improvement.
- Scaler experiment: MinMaxScaler substantially outperformed RobustScaler for LSTM specifically (MAE 26.15 vs 28.91); StandardScaler caused severe instability ($R^2 = -0.99$).
- Bidirectional LSTM test: a combined search (`window`/`bidirectional`/`units`/`lr`) selected a non-bidirectional configuration as best, indicating bidirectionality did not offer a clear advantage for this dataset.

Final LSTM (MinMaxScaler, `window`=12, 64→32 units, seeded for reproducibility): MAE=26.17, RMSE=62.52, MAPE=22.06%, R^2 =0.493 - the best **LSTM result across all versions tested, finally exceeding the Naive baseline.**

Reproducibility note: LSTM results showed run-to-run variance (MAE ranging 26.15–28.56) attributable to neural network weight initialization randomness. A fixed random seed (42) was applied to the final reported configuration to ensure a reproducible result.

Cross-cutting finding

Across both XGBoost (v2 regularization) and LSTM (v2 added complexity, dropout ablation, and scaler testing), added model complexity consistently failed to improve - and sometimes worsened performance. This reinforced a theme observed throughout the project: results here are bounded more by the dataset's scale and inherent structure than by architectural sophistication.

8. Challenges and Solutions

Box-Cox inversion instability: with a negative `lambda` (-0.531), the inverse transform (`y ** (1/lambda)`) is undefined for negative bases and highly sensitive near the domain boundary. Unconstrained models (particularly Linear Regression) occasionally produced predictions triggering NaN or astronomically large values upon inversion. Resolved via a safe inversion function with domain clipping and a sanity-bounded maximum output.

Combining Box-Cox with the train/test ratio experiment initially produced extreme, uninterpretable MAE values (billions) due to this same instability compounding across larger test sets. Resolved by isolating the ratio experiment to the raw target, keeping the Box-Cox decision separately validated.

Testing whether engineered lag/rolling features were redundant for LSTM (hypothesis: a sequence model should learn temporal patterns natively from raw history alone): disproven empirically, removing these features caused severe performance collapse (MAE 29→57, R^2 0.42→-0.26), showing engineered features remain valuable even for models designed to process sequences directly, likely due to the modest network size and training data volume.

LSTM result non-determinism: addressed via fixed random seeding for the final reported model, after identifying and diagnosing run-to-run variance across repeated identical-code executions.

Lag features computed after train/test split initially discarded a small number of valid boundary rows (~10 rows, <0.3% of test data) - acknowledged as a minor, documented trade-off rather than requiring a full pipeline rebuild given the negligible practical impact.

9. Conclusion

XGBoost (tuned) is the recommended final model, achieving MAE=17.91 Wh, RMSE=44.86 Wh, MAPE=14.61%, and $R^2=0.744$ - the strongest result across every baseline and deep learning architecture tested. This outperforms even the benchmark literature's best reported figure on this identical dataset (GRU, $R^2=0.62$), attributable largely to the rolling-average and lag features engineered here, which alone accounted for approximately 75% of the model's predictive power.

Deep learning models (LSTM, GRU, CNN-GRU hybrid) were designed, tuned, and rigorously compared as required, with the best deep learning result (LSTM, $R^2=0.493$) ultimately falling short of XGBoost. This is a well-documented and empirically confirmed pattern for tabular, moderately sized time-series datasets (~16k rows), where tree-based methods frequently outperform deep architectures that typically require substantially larger data volumes to realize their representational advantages.

Every major methodological decision in this project - outlier treatment, target transformation, scaling method, split ratio, sequence window, optimizer, dropout, and hyperparameters - was tested empirically against alternatives rather than assumed, with each decision traceable back to concrete evidence from EDA or dedicated experimentation. A recurring theme throughout was that simpler configurations consistently matched or outperformed more complex alternatives, reflecting the dataset's inherent scale and structure rather than a limitation of experimentation depth.

Future work: quantile regression or a specialized high-usage-tail model to address the identified spike-underestimation limitation; attention-based sequence architectures, should additional training data become available; extending the holiday/domain-specific feature set with richer contextual signals (e.g., occupancy schedules, if available).

10. References

1. Candanedo, L. M., Feldheim, V., & Deramaix, D. (2017). Data driven prediction models of energy use of appliances in a low-energy house. *Energy and Buildings*, 140, 81–97. <https://doi.org/10.1016/j.enbuild.2017.01.083>
2. Al-Obeidat, F., Tubaishat, A., Shah, B., & Delen, D. (2022). Deep learning-based sequence modeling for residential energy consumption prediction using multivariate time-series data. *Journal of Big Data*, 9(1), 45–68. <https://doi.org/10.1186/s40537-022-00591-2>
3. Pan, H., Tang, X., Wang, J., & Zhou, B. (2020). A hybrid CNN-GRU deep learning strategy for residential energy prediction. *Energies*, 13(12), 3241. <https://doi.org/10.3390/en13123241>